DEVELOPMENT
DATA PARTNERSHIP

Unlocking Data Silos for AI Applications

Supported by: Gates Foundation

Program Objectives

A scalable, rights-based approach to unlocking proprietary datasets for building public-good AI applications

This Gates Foundation–funded initiative aims to **develop a scalable, replicable model for unlocking proprietary data for use in AI applications.** The first part supports a framework for building national low-resource language libraries derived from media company and government content, and the second supports facilitating scaled access to donor grantee datasets.

Current efforts to gather low-resource language data rely on scraped web content and isolated field recordings. While useful, these sources are fragmented and insufficient for training high-performing models. As a result, speakers of these languages are often left behind. This project addresses the gap by using the Development Data Partnership's infrastructure to license and prepare high-value content—such as news, books, radio, and surveys—under clear consent, privacy protections, and non-commercial terms. The resulting rights-based libraries will be transferred to national stewards to support continued corpus growth, AI training, and inclusive digital transformation rooted in local languages and priorities.

Similarly, donor grantee datasets are often siloed under varying legal terms and organizational policies. This initiative seeks to harmonize governance and technical frameworks to streamline responsible access

DEVELOPMENT DATA PARTNERSHIP

The Development Data Partnership

Introduction

Why
Data Partnership?

- Public policy and provision of public infrastructure and services are heavily dependent on data – higher quality, timely data translates into more effective sector and program prioritization, design, implementation, monitoring, and evaluation.

- Increasingly, the private sector is generating data that could be used to complement traditional public-sector data collection methods. Through public-private collaboration, in addition to generating more timely and relevant data for decision making, entirely new public good use cases could be discovered and implemented.

DEVELOPMENT
DATA PARTNERSHIP

Development Data Partnership Goals

- Coordinate and aggregate private sector data demand on behalf of public actors
- Reduce international organization duplication of effort and facilitate collaboration on solutions
- Link public sector challenges and domain expertise to relevant proprietary data
- Significantly reduce the transaction costs associated with data sharing
- Increase transparency and accountability for integrating use of proprietary data in public good analytics
- Increase the capacity of the public sector to procure and use proprietary data and platforms for better decision making

DEVELOPMENT DATA PARTNERSHIP

Our Partners | Companies

Our Partners | International Organizations

AFRICAN DEVELOPMENT BANK GROUP

ADB

European Bank
for Reconstruction and Development

European Investment Bank

IDB | Invest

IDB

INTERNATIONAL MONETARY FUND

OECD

unicef
for every child

UNDP

THE WORLD BANK
IBRD • IDA | WORLD BANK GROUP

DEVELOPMENT
DATA PARTNERSHIP

Partnership Frameworks

The Partnership supports international development through data sharing and data science collaboration between companies and international organizations, leveraging three frameworks:

Legal

Governance

Technical

DEVELOPMENT DATA PARTNERSHIP

How The Partnership Works

1

Data Partnership

Data License Agreements are signed between the data partner and each participating organizations (e.g. World Bank, IMF, IDB, OECD,UNDP)

2

Project Proposals

Data partner uses the Partnership Portal to evaluate all proposals for using their data in a project.

3

Data Management

Upon project approval, data are securely and responsibly managed on behalf of all staff through the Partnership IT architecture and procedures.

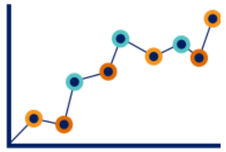

4

Data Goods

Data partners are kept updated through the Partnership Portal on derived data products and produced code.

DEVELOPMENT DATA PARTNERSHIP

The Partnership Today

Scalable Partnership Model

Each year, international organization participation continues to grow, with 11 partners today.

Over 400 Development Projects

Teams from across membership organizations are leveraging data provided by the private sector to address global challenges and the Sustainable Development Goals.

Growing Data Partners

30+ active data-sharing agreements signed with private companies to bridge private sector data and public sector needs.

Impactful Results

Insights and data goods generated by the projects are informing public policy and government decision-making across different countries and regions.

DEVELOPMENT
DATA PARTNERSHIP

Learn About Our Projects

Leveraging LinkedIn Data to Understand the Green and Digital Transformations in the Labor Market and the Future of Work

Case Study | LinkedIn | Jobs and Development

The economy and the labor market have evolved due to digital and green transitions. New technologies and new forms of work stemming from digitalization, climate change, and ongoing efforts to shift to a low-carbon economy have been changing jobs and skills at a rapid pace.

Learn More

How to Leverage Data for Better Transport, Digital Connectivity, and Sustainable Development in the Amazon

Case Study | Google | Mapbox | Meta | OOkla | Digital Development | Transport

The World Bank Transport team used data from Google, Mapbox, Meta, and Ookla to identify infrastructure gaps in the Amazon areas of Brazil, Colombia and Peru.

Learn More

Can private management of African protected areas improve socioeconomic and wildlife outcomes?

Case Study | AtlasAI | Inequality and Shared Prosperity

African governments increasingly entrust the management of protected areas to private, nongovernmental organizations, with the hope that the significant resources and technical capacities of private organizations will help realize the potential of these areas.

Learn More

Advancing Women's Empowerment Through Data

Case Study

Data plays a fundamental role in identifying the challenges women are facing and in allowing policymakers and businesses to make informed decisions that improve women's well-being and access to economic opportunities and overall development outcomes.

Learn More

Unleashing the Power of Data to Tackle Traffic Congestion and Promote Road Safety

Case Study | Mapbox | Mapillary | Waze

As the global population is constantly increasing, building modern and sustainable cities to accommodate everyone is crucial. This is enshrined in the 2030 Agenda for Sustainable Development, in which the Sustainable Development Goal 11 aims to make cities more inclusive, safe, resilient and sustainable.

Learn More

Leveraging Social Media Data to Map Road Traffic Crashes

Case Study | X

Challenge Road traffic crashes are among the world's most pressing public health challenges. Crashes are the leading cause of death for those 5-29 years old and are the 8th leading cause of death considering all ages.

Learn More

Impact Stories

DEVELOPMENT DATA PARTNERSHIP

Component 1

Unlocking Low-Resource Language Data

DEVELOPMENT DATA PARTNERSHIP

Unlocking Low-Resource Language Data Silos

Component 1 | Low-Resource Language Data

Component Objective
To create a scalable, replicable model for building national digital language libraries for training and tuning AI models.

Challenge
~4 billion people speak low-resource languages, which are often overlooked in AI model training and applications. Widely spoken low-resource language content can exist but is often proprietary and locked in silos.

Solution
Leverage Partnership legal, governance, and technical frameworks to unlock data from media companies and government agencies for use in public good AI model training and applications.

DEVELOPMENT DATA PARTNERSHIP

Why Secure Licensing to Complement Open Data Initiatives

Ethical Compensation Framework | Current AI training practices create an unfair extraction model where creators provide value without receiving compensation. A secure licensing system establishes fair value exchange where all creators are compensated proportionally to their contribution, regardless of geography.

Data Scarcity Reality | Most low-resource languages have less than 1% of the data available for English models. Open data mandates cannot bridge this gap because the high-quality content doesn't exist in open repositories, and manual collection efforts can be challenging to scale and sustain.

Quality and Safety Benefits | Copyrighted content offers advantages through professional editorial processes and fact-checking that scraped data can lack. This results in higher linguistic accuracy. Using professionally curated content reduces the substantial costs of bias mitigation and content moderation required when scraping unfiltered internet data.

Economic Sustainability | This system creates sustainable revenue streams for publishers and content creators while enabling breakthrough research. Market mechanisms drive investment in underrepresented languages more effectively than charity models, creating a self-sustaining ecosystem where quality content fuels quality research and linguistic diversity thrives.

'Impossible' to create AI tools like ChatGPT without copyrighted material, OpenAI says

Pressure grows on artificial intelligence firms over the content used to train their products

https://www.theguardian.com/technology/2024/jan/08/ai-tools-chatgpt-copyrighted-material-openai

OpenAI Strikes a Deal to License News Corp Content

The deal gives OpenAI's chatbots access to new and archived material from The Wall Street Journal, The New York Post, MarketWatch and Barron's, among others.

"...the Wall Street Journal reported the agreement could be worth as much as $250 million over five years."

https://www.nytimes.com/2024/05/22/business/media/openai-news-corp-content-deal.html

DEVELOPMENT DATA PARTNERSHIP

Approach: Learning by Doing

Component 1 | Low-Resource Language Data

1

Legal Framework | Curation

The team will engage and source content from media companies and government agencies in a pilot country, leveraging a modified version of the DDP master data license agreement. The team will set up a data review and cleaning pipeline and meta data schema tailored for proprietary low-resource language libraries.

2

Technical Framework | Pre-Processing

The team will develop, document, and implement methods for pre-processing text and audio data received from different sources -- .pdf newspaper scans, household survey audio recordings, .xml newspaper web pages, etc. The team will create a data flow that ensures original IP are not fully disclosed through the library, including methods for chunking and shuffling data, linking chunks to context, and making them accessible via a monitored API.

3

Technical Framework | Library Management

The team will leverage the current Partnership processes for managing ad fulfilling data requests, with the addition of a new externally facing application page and catalog. The team will prepare a model verification tool to ensure original proprietary contents cannot be ‘leaked’ through the model, as well as other guidance and benchmarking tools for users.

4

Governance Framework | Scaling & Sustaining

Finally, the team will support knowledge transfer to the pilot country government for replicating the library establishment and process, supported by a percentage of funds from third party companies purchasing updates of the library contents. All outputs will be packaged for use in new countries just starting their libraries.

DEVELOPMENT DATA PARTNERSHIP

Use Case: Malawi Implementation

Component 1 | Low-Resource Language Data

- Low Resource Language
  - Chichewa spoken by 20+ million people (of which ~25% speak English)

- Government Partners
  - Ministry of Information:
    - Key national language library stakeholder
    - Will use project-derived LLMs to support government e-services program
  - National Statistical Office
    - Providing household survey audio recordings and transcriptions
    - Will use project-derived ASR model to support survey supervision

- Media Partners
  - Organized workshop with 27 media, academic, government, and NGO stakeholders
  - Bi-lateral negotiations with 18 companies, using team-developed legal terms and cost recovery templates

DEVELOPMENT DATA PARTNERSHIP

Curation

Stakeholder Engagement

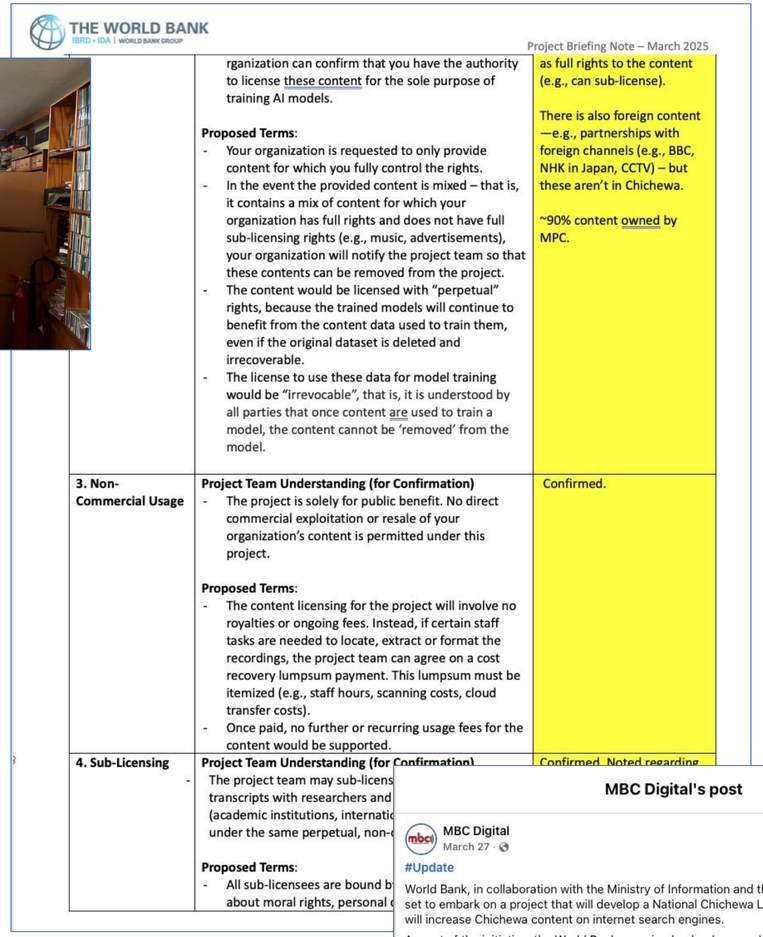

THE WORLD BANK
IBRD • IDA | WORLD BANK GROUP

Project Briefing Note – March 2025

	rganization can confirm that you have the authority to license these content for the sole purpose of training AI models. **Proposed Terms**: - Your organization is requested to only provide content for which you fully control the rights. - In the event the provided content is mixed – that is, it contains a mix of content for which your organization has full rights and does not have full sub-licensing rights (e.g., music, advertisements), your organization will notify the project team so that these contents can be removed from the project. - The content would be licensed with "perpetual" rights, because the trained models will continue to benefit from the content data used to train them, even if the original dataset is deleted and irrecoverable. - The license to use these data for model training would be "irrevocable", that is, it is understood by all parties that once content are used to train a model, the content cannot be 'removed' from the model.	as full rights to the content (e.g., can sub-license). There is also foreign content —e.g., partnerships with foreign channels (e.g., BBC, NHK in Japan, CCTV) – but these aren't in Chichewa. ~90% content owned by MPC.
3. Non-Commercial Usage	**Project Team Understanding (for Confirmation)** - The project is solely for public benefit. No direct commercial exploitation or resale of your organization's content is permitted under this project. **Proposed Terms**: - The content licensing for the project will involve no royalties or ongoing fees. Instead, if certain staff tasks are needed to locate, extract or format the recordings, the project team can agree on a cost recovery lumpsum payment. This lumpsum must be itemized (e.g., staff hours, scanning costs, cloud transfer costs). - Once paid, no further or recurring usage fees for the content would be supported.	Confirmed.
4. Sub-Licensing	**Project Team Understanding (for Confirmation)** - The project team may sub-licens transcripts with researchers and (academic institutions, internati under the same perpetual, non- **Proposed Terms:** - All sub-licensees are bound b about moral rights, personal	Confirmed. Noted regarding

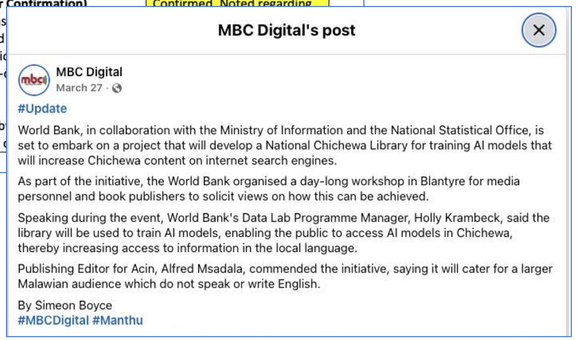

MBC Digital's post

MBC Digital
March 27 ·

#Update

World Bank, in collaboration with the Ministry of Information and the National Statistical Office, is set to embark on a project that will develop a National Chichewa Library for training AI models that will increase Chichewa content on internet search engines.

As part of the initiative, the World Bank organised a day-long workshop in Blantyre for media personnel and book publishers to solicit views on how this can be achieved.

Speaking during the event, World Bank's Data Lab Programme Manager, Holly Krambeck, said the library will be used to train AI models, enabling the public to access AI models in Chichewa, thereby increasing access to information in the local language.

Publishing Editor for Acin, Alfred Msadala, commended the initiative, saying it will cater for a larger Malawian audience which do not speak or write English.

By Simeon Boyce
#MBCDigital #Manthu

Objective: Efficiently engage with a wide range of content stakeholders, ensuring transparent communication about data provenance and planned usage.

Approach: Kick-off by organizing a publicly communicated and government-supported stakeholder workshop with content providers for awareness and inputs. Engage in follow-up negotiations using a "terms sheet" for efficient and transparent communication about data license agreement contents.

In Practice: Media stakeholder workshop held in Malawi in 3/2025, with 27 media stakeholders – newspapers, radio, book publishers, writers' union, copyright association, and universities – as well as senior government leadership from the Ministry of Information and the National Statistical Office. Workshop followed by 18 bi-lateral stakeholder meetings using the **terms sheet**, securing buy-in for content provision.

Curation

Legal Framework

Objective: Secure rights to proprietary content from media firms for distribution through a library and for use in training AI models, while minimizing negotiation requirements and protecting both content providers and users.

Approach: Leverage Development Data Partnership Master Data License Agreement (MDLA), which has been signed by more than 40 international organizations and private companies for use in sharing sensitive and proprietary data. Supplement with a new AI Annex, which grants sub-licensing rights, permissions for AI model training, and protections for content providers and users.

In Practice: To ensure fairness and scalability, the new MDLA content has been reviewed and revised pro-bono by Freshfields, a global-leading law firm.

Freshfields Bruckhaus Deringer

Freshfields Bruckhaus Deringer is one of the world's leading law firms, part of London's elite "Magic Circle," with over 275 years of experience advising on complex global deals and disputes.

Their decision to provide **pro bono support** on our AI content licensing agreement brings **world-class legal expertise—normally reserved for high-stakes corporate clients**—directly to our initiative. This **underscores the significance of our project and the trust placed in its global relevance.**

DEVELOPMENT DATA PARTNERSHIP

Curation

Meta Data Schema

Objective: Make data as discoverable and usable as possible for AI model training and fine tuning.

Approach: Leverage the World Bank's open-source Meta Data Editor to create template for an updated meta data schema that meets the unique needs of a text and audio content library for model training.

In Practice: The team is currently developing the schema and template. The team will work with content providers to ensure templates are completed. Over time, some components may become automatable (e.g., context summaries).

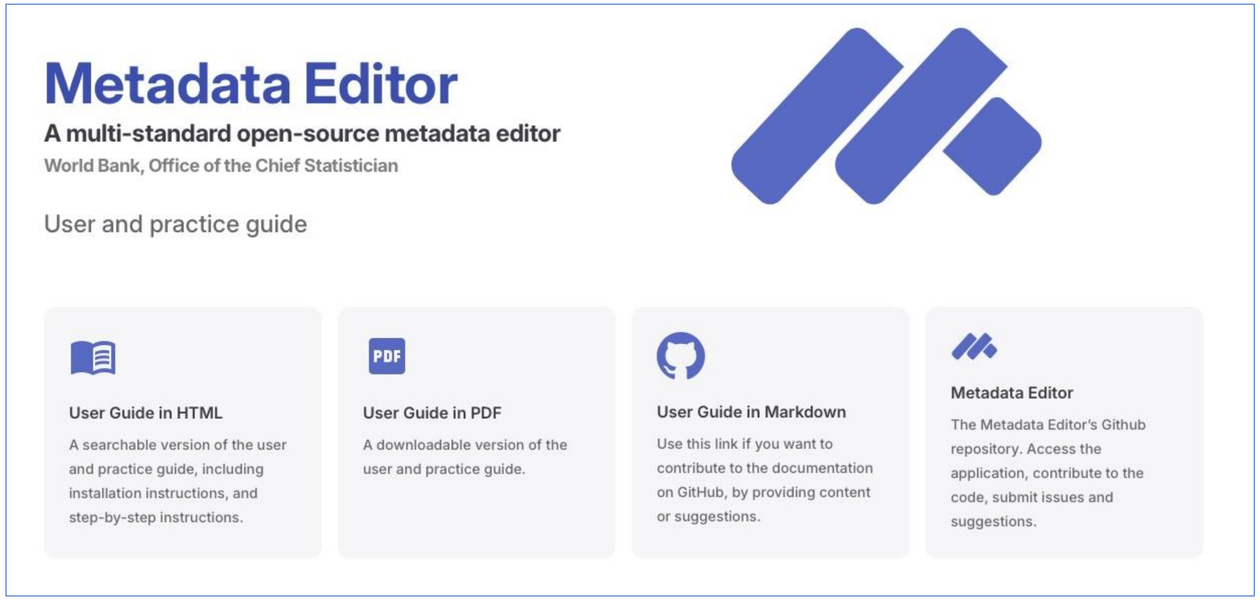

DEVELOPMENT DATA PARTNERSHIP

Pre-Processing

Making Data AI-Model Ready

Objective: Reduce duplication of effort in making data AI-model ready, while providing the research community with tools needed to scale pre-processing work over time.

Approach: First, develop method to transform received formats into .txt and .mp3 files. Then, generate statistics on the datasets, to identify common themes, potential errors, biases, etc. Finally, clean the data by removing duplications, unnecessary formatting (e.g., xml notation, PII, data we do not have warrant to share, etc.) with as much automation as possible .

In Practice: Leveraging 5,000 scanned Chichewa newspaper pages and working with local students, build Chichewa OCR model for automating conversion of newsprint pdf to text. Working with the National Statistical Office, developing method to anonymize recorded household surveys.

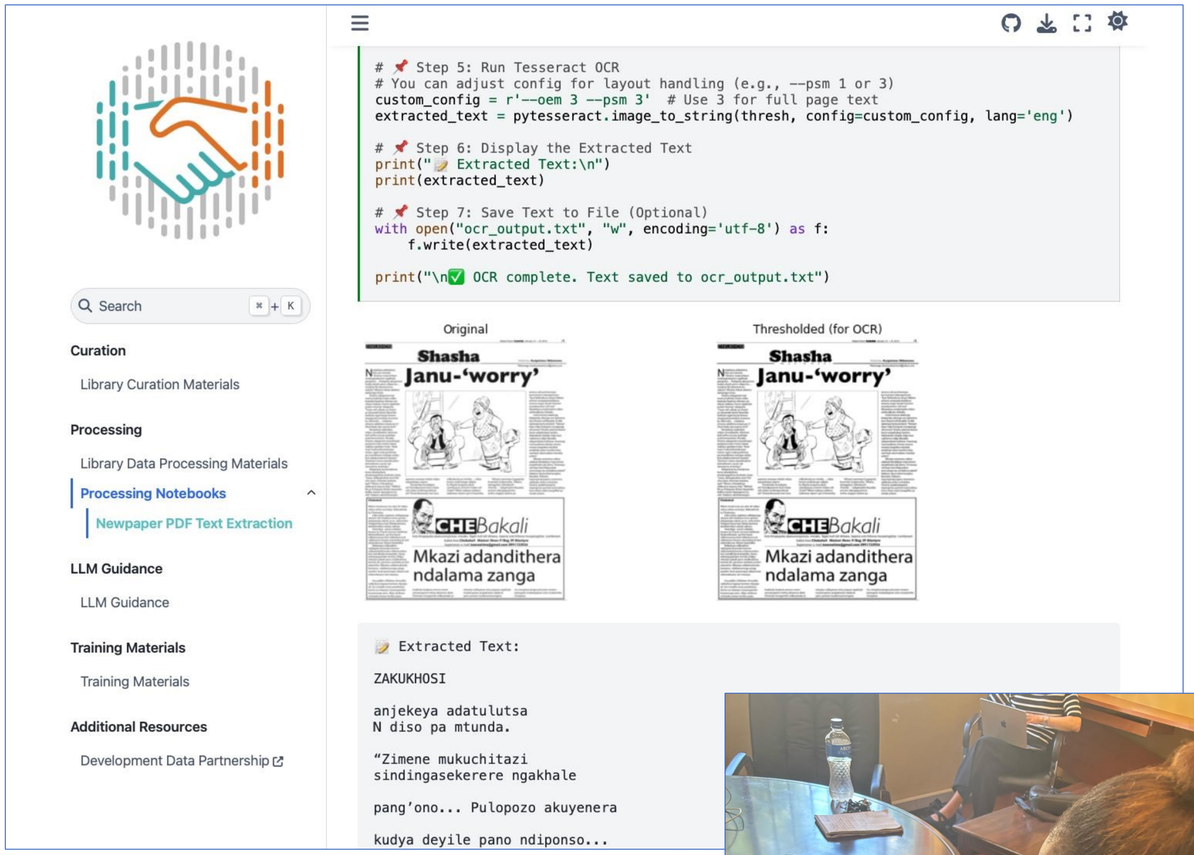

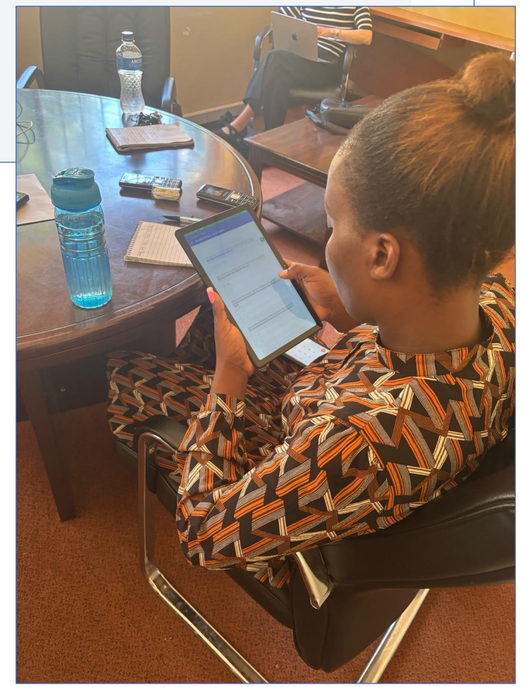

Pre-Processing

Protecting IP Rights

Objective: Mitigate risk of non-compliance with license agreements by limiting ability to download complete original copyright contents, without need to set up expensive and cumbersome self-contained trusted compute environments.

Approach: To prevent copyright infringement, chunk data based on section headers and file length (e.g., for text, page-section-paragraph-sentence groups) and shuffle. Add context summary to the meta data for each chunk, which will be callable via API. The API is being designed to include limiters and monitoring tools, to ensure original contents are not stitched together by users, as well as injection of 'canaries' that are assigned to each API-registrant, to support future license agreement enforcement.

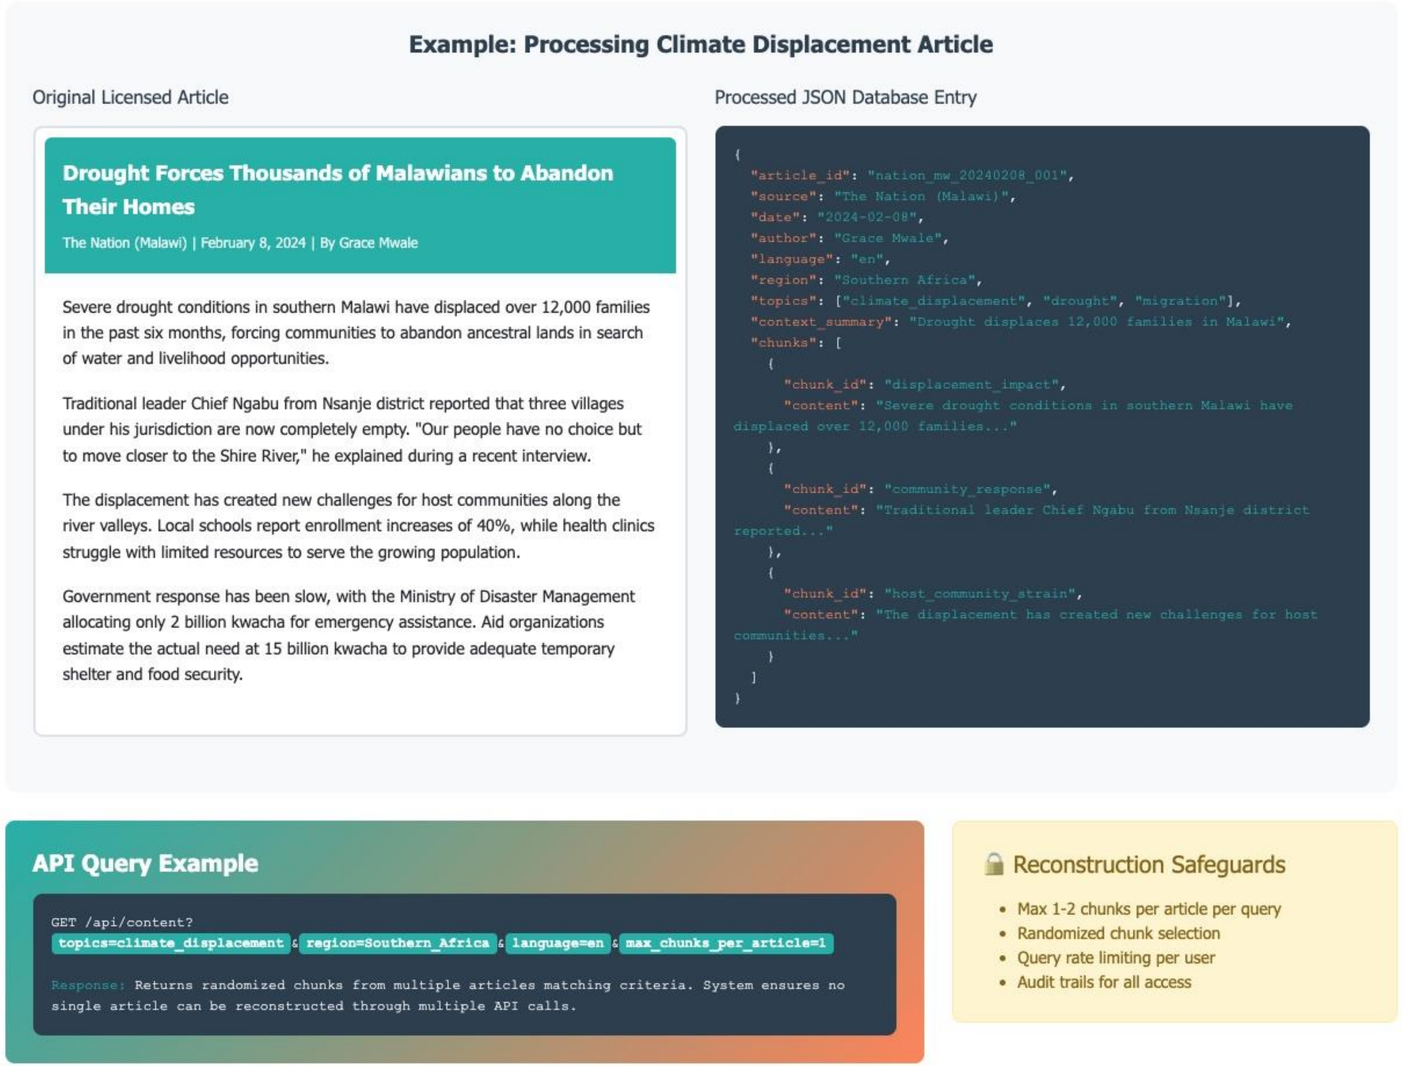

Screenshot: API and IP Protection Concept (Claude.AI)

DEVELOPMENT DATA PARTNERSHIP

Library Management

Supporting Partnership Members and Data Fellows

Objective: Make access to the library seamless for partner international organization staff and Data Fellows

Approach: On the back end, leverage the Development Data Partnership's existing framework for receiving and reviewing proposals, securing agreement to data licensing terms, provisioning data access, providing supporting documentation and technical assistance, and monitoring usage.

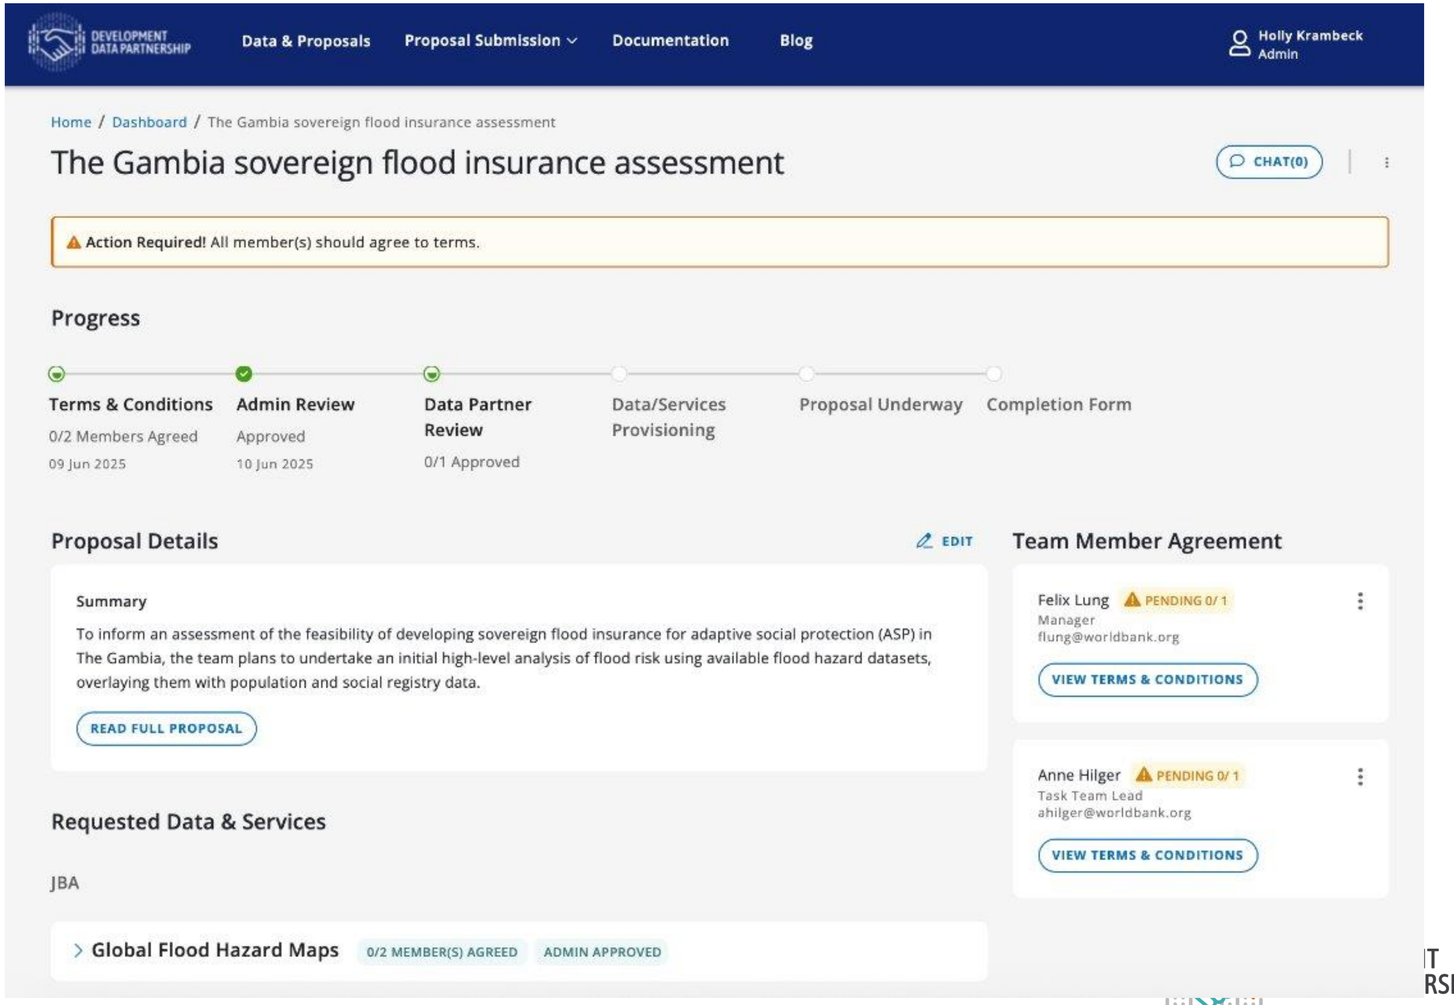

Screenshot: Development Data Partnership Project Management Portal

Library Management

Supporting Scaling

Objective: Define 1-3 solutions from different providers (AWS, Google Cloud, Azure) for scaling national libraries.

Approach: On the back end, prototype purpose-built data governance solutions for managing, fulfilling, and monitoring requests, leveraging support from technology partners. On the front end, create a public-facing application and meta-data catalog for researchers seeking use of the library.

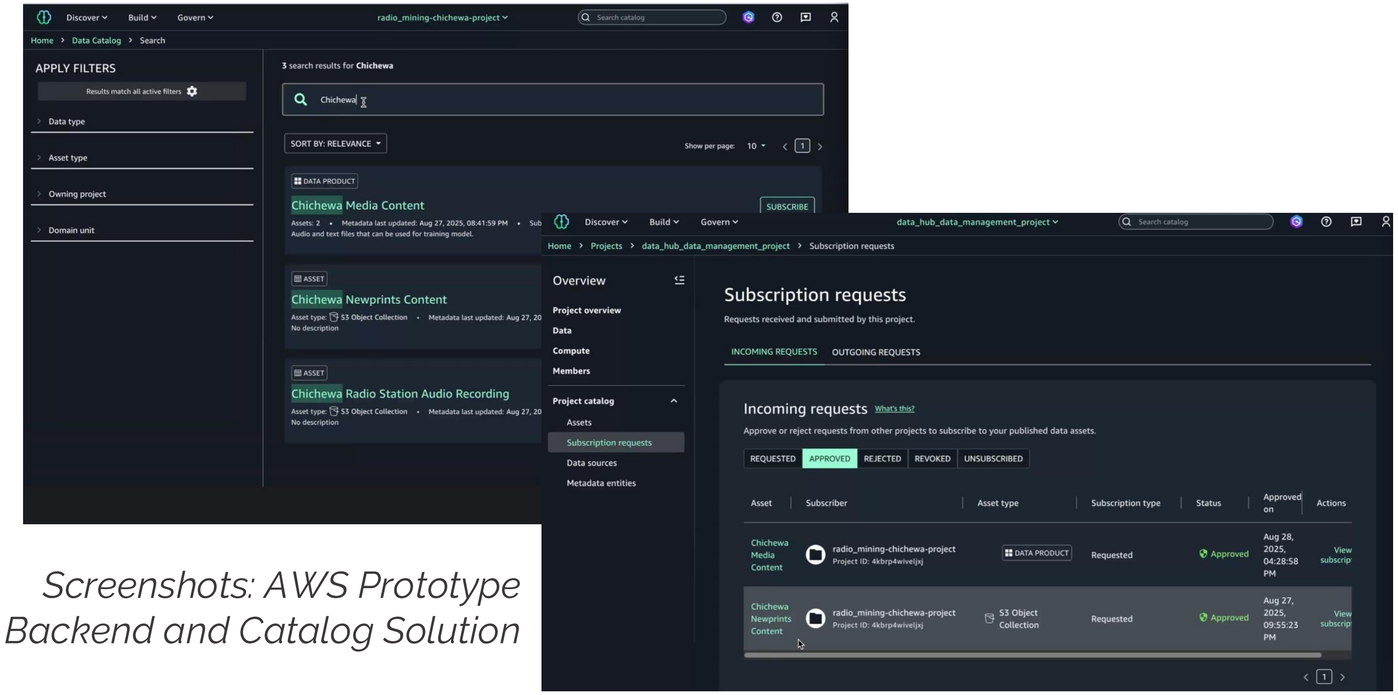

Screenshots: AWS Prototype Backend and Catalog Solution

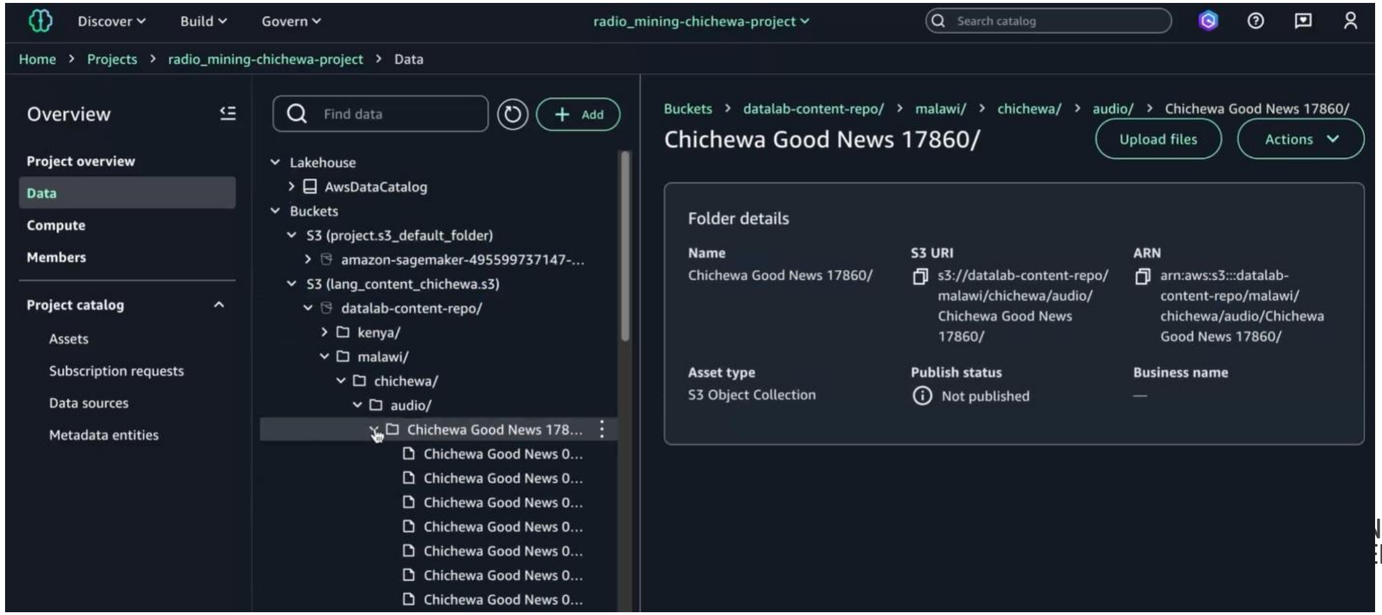

Library Management

Guidance and Tools for Users

Objective: Mitigate risk of non-compliance with license agreements and reduce duplication of effort.

Approach: As part of the sub-licensing agreement, require researchers to submit results of a model test provided through the Partnesrhip to ensure fine-tuned models minimize 'leakage' of original IP. In addition, provide researchers with other accompanying documentation in the Partnership's living repository of code examples and guidance.

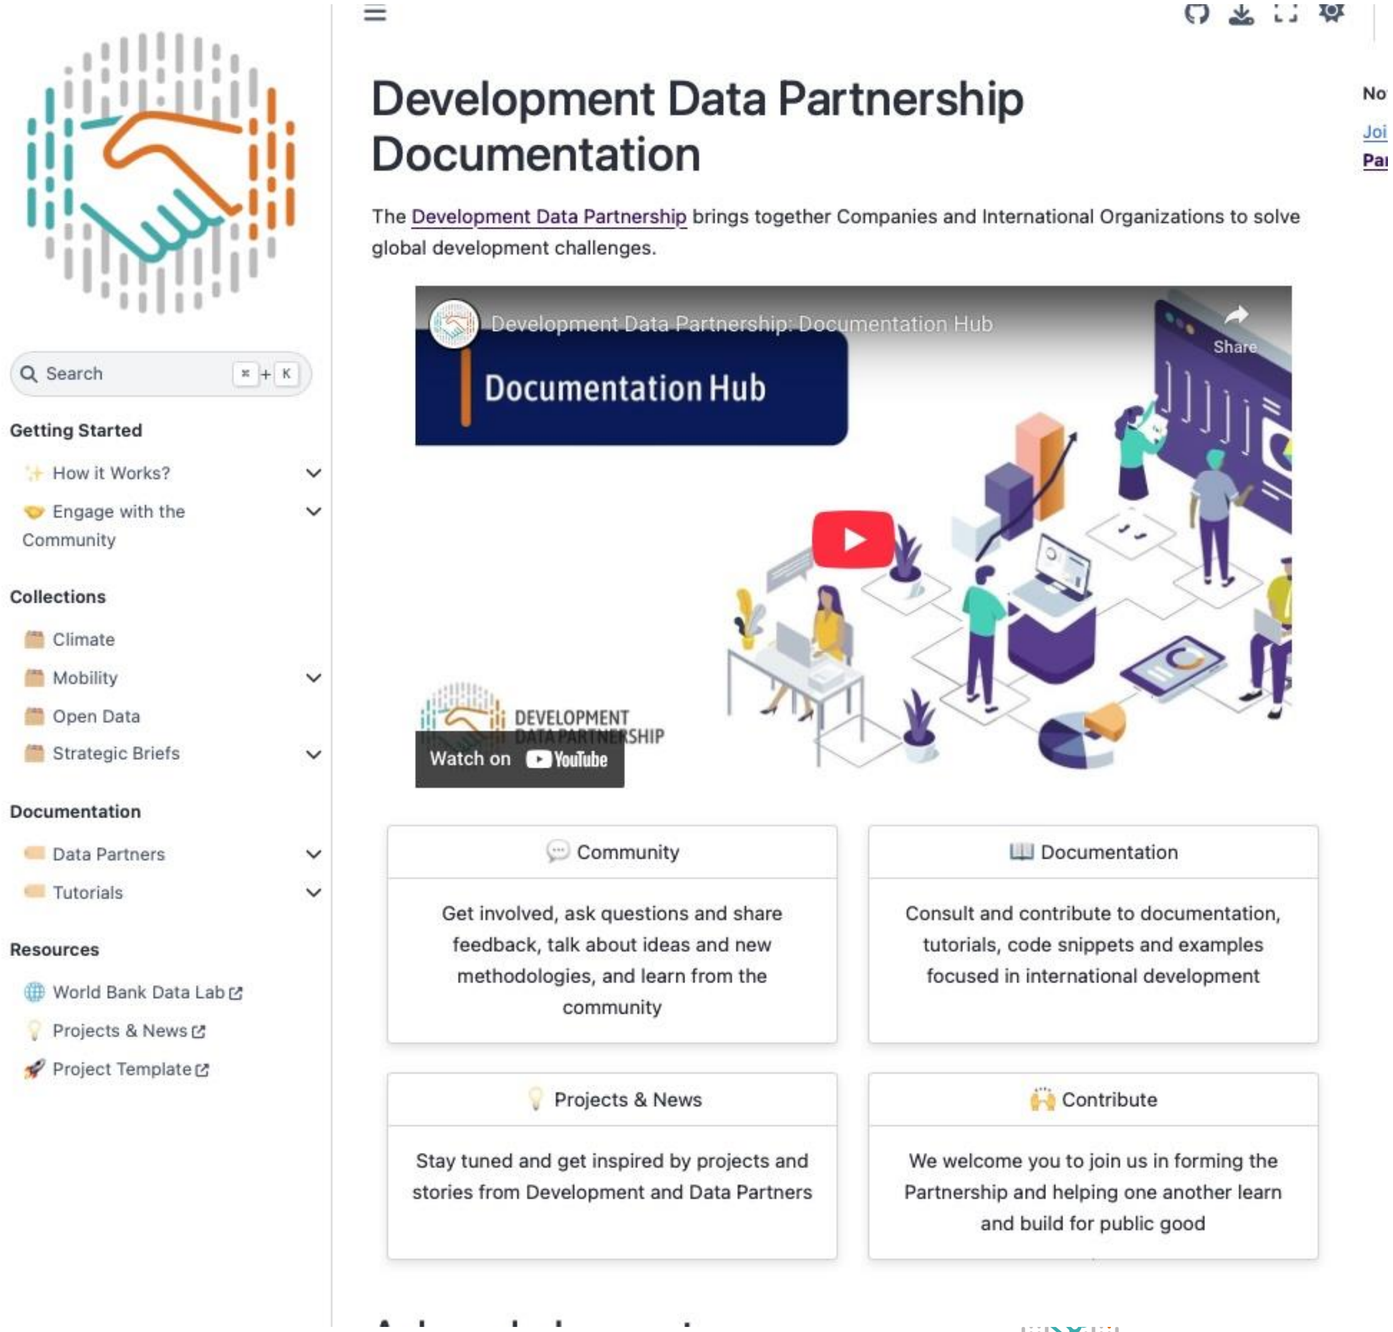

Screenshot: Development Data Partnership Documentation Portal

Scaling and Sustaining

Transferring Management to Country

Objective: Transfer ongoing stewardship to the national government for sustained content additions and revenue generation.

Approach: Learning by doing, through the Malawi pilot, work with Ministry of Information, University of Malawi, and MAREN (Malawi Research and Education Network) to establish a sustaining stewardship and funding model for library content provisioning, with different API tiers for research and commercial use, as well as a built-in sustained update and funding mechanism.

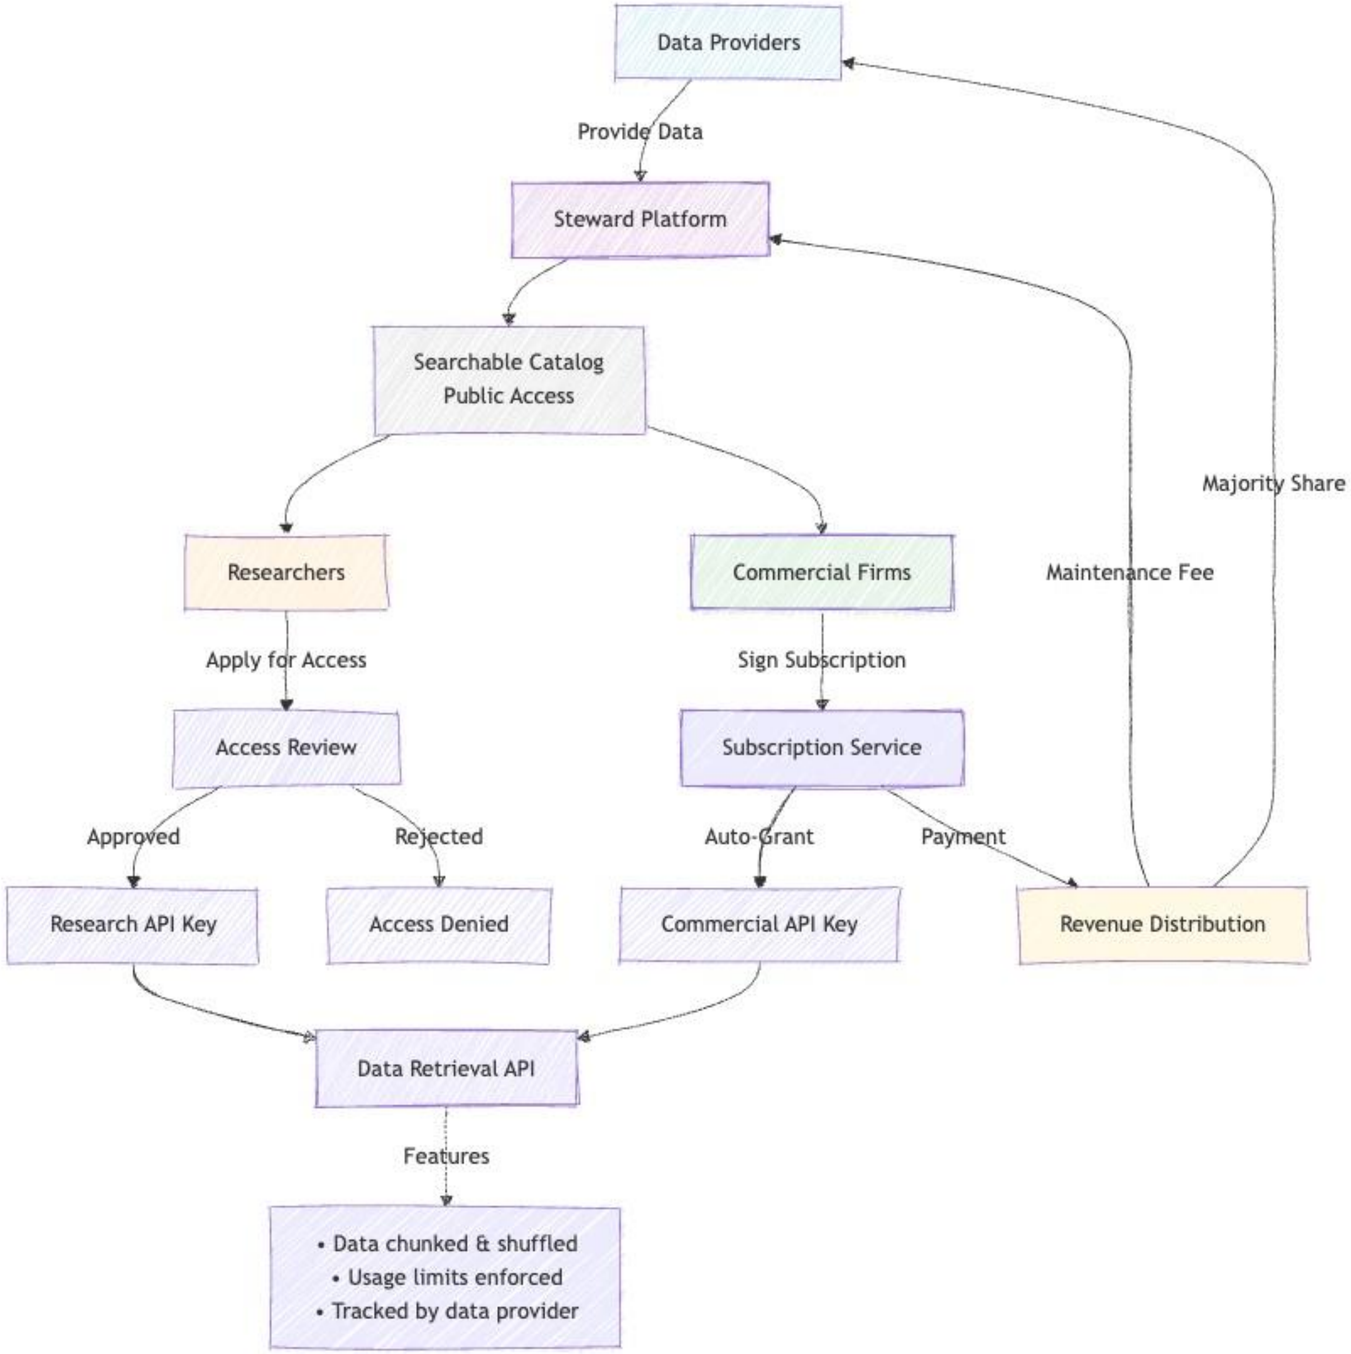

Scaling and Sustaining

Replicating Model with Other Governments

Objective: Build capacity of other governments to build their own libraries.

Approach: Prepare templates, guidance, and code accessible through a single open GitHub repository via living and dynamically-generated web book. Through the World Bank Digital Transformation investment team, the Partnership team has already received requests to support and fund replication of the Malawi program in Zambia and Uzbekistan, demonstrating the strong demand for these models and frameworks.

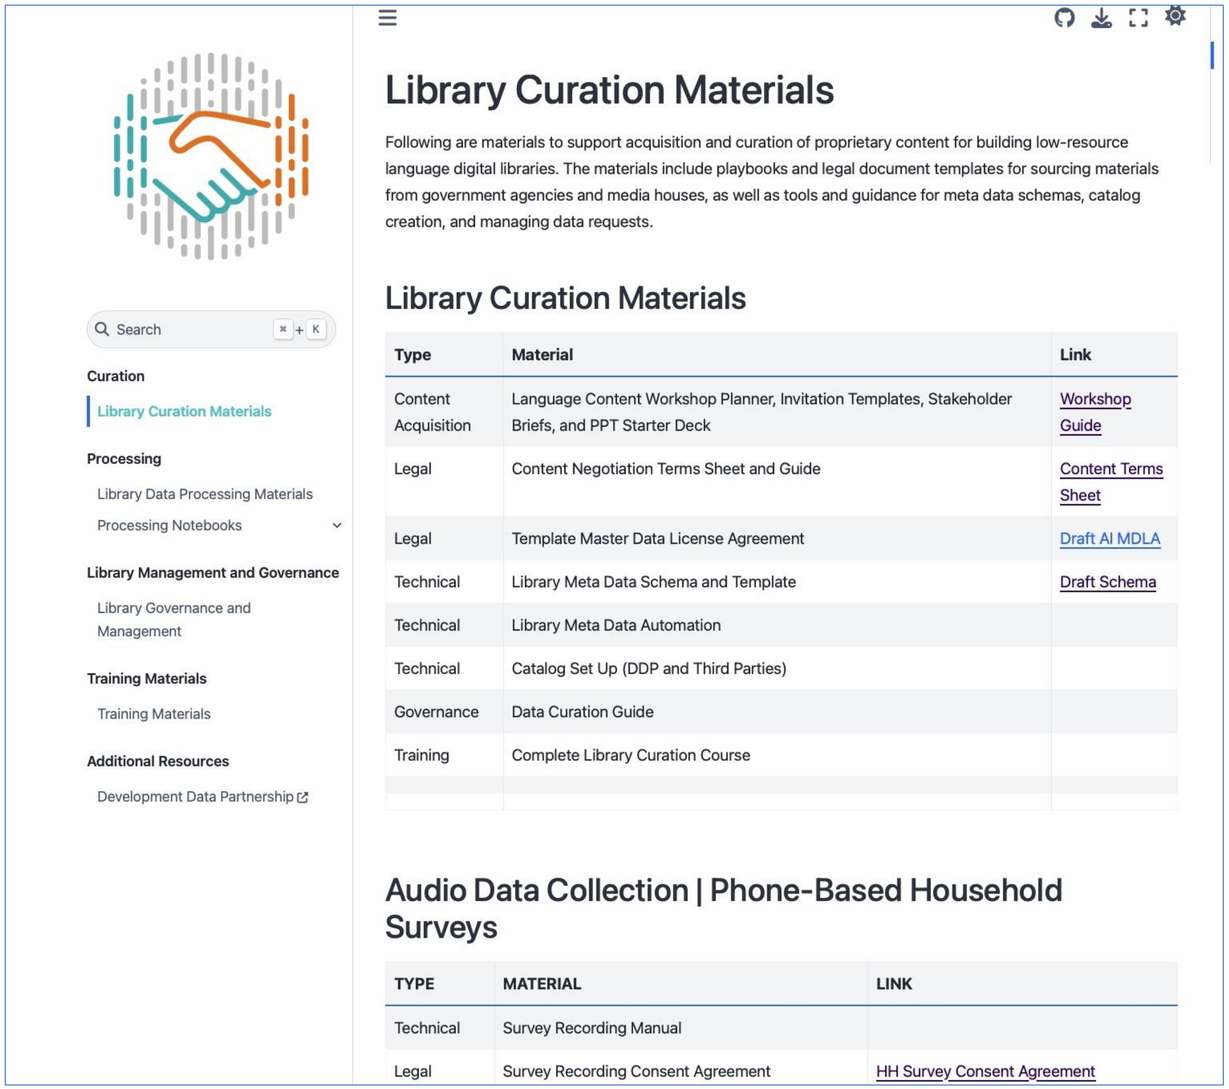

https://bit.ly/data-for-AI-library

Scaling and Sustaining

Replicating Model with Other Governments

Objective: Build capacity of other governments to build their own libraries.

Approach: Prepare templates, guidance, and code accessible through a single open GitHub repository via living and dynamically-generated web book. Through the World Bank Digital Transformation investment team, the Partnership team has already received requests to support and fund replication of the Malawi program in Zambia and Uzbekistan, demonstrating the strong demand for these models and frameworks.

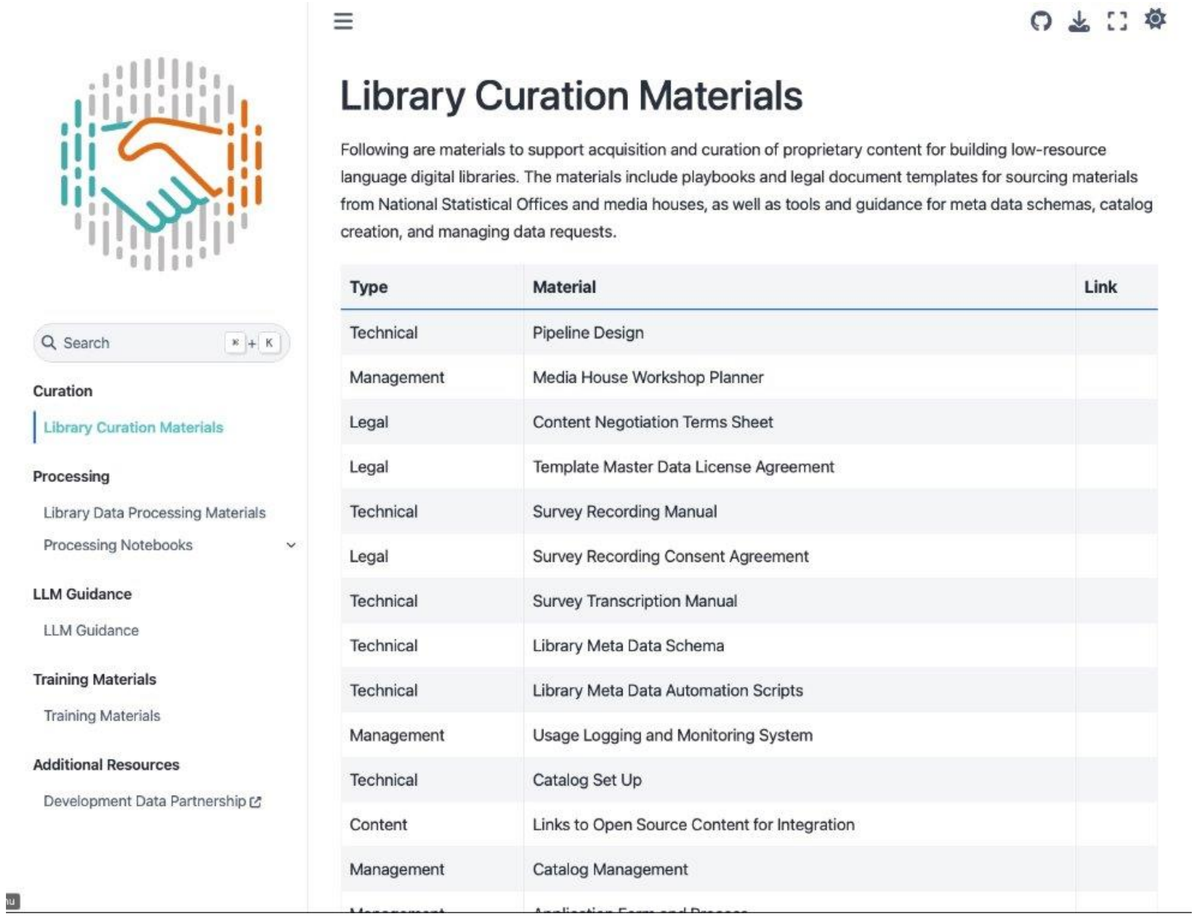

https://bit.ly/data-for-AI-library

DEVELOPMENT DATA PARTNERSHIP

Component 2

Unlocking Grantee Datasets

Unlocking Grantee Data

Component 2 | Grantee Data

Component Objective
To create a scalable, replicable model for provisioning access to non-open grantee datasets for public good.

Challenge
The Gates Foundation funds dozens of organizations each year to support data collection, and while dissemination is a requirement, grantees often struggle with the responsibility of data provisioning months and years after the grant has closed. With no unified catalog, legal agreement, or data architecture, third parties are challenged to discover or access these potentially useful datasets.

Solution
Leverage Partnership legal, governance, and technical frameworks to unlock data from Gates Foundation grantees for use in public good research and AI applications.

DEVELOPMENT DATA PARTNERSHIP

Approach: Learning by Doing

Component 2 | Grantee Data

1

Legal Framework | Curation

The team will engage and source content from Gates Foundation grantees, leveraging a modified version of the DDP master data license agreement. The team will leverage the existing DDP infrastructure to ingest, store, and provision datasets.

2

Technical Framework | Pre-Processing

As is needed, the team may engage in pre-processing of datasets to support ease of use, as well as preparation of APIs and/or documentation.

3

Technical Framework | Library Management

The team will adjust the DDP data application and provisioning process for the project. The team will prepare third-party engagement and data provisioning workflow that is more easily scalable than the current model.

4

Governance Framework | Scaling & Sustaining

Oversight will be managed through the Partnership's current inclusive governance framework. The team will support an outreach effort across the international organization partners for awareness building and training on using the hosted datasets and may also support a kick-off hack-a-thon style event to encourage more third-party researchers to apply for and use the datasets.

DEVELOPMENT DATA PARTNERSHIP

Project Timeline

Key Milestones

Project Management

How We Organize Our Work

The core team includes the Development Data Partnership team, World Bank legal counsel and data scientists, and consultants with specialized experience in working in the Malawi context and with low-resource languages. The extended team include colleagues from the World Bank IT department, as well as the Development Economics Data Group and Digital Transformation teams.

Project tasks are organized and monitored using a GitHub task tracker, visible to team members for transparant project management.

Project GitHub Task Tracker

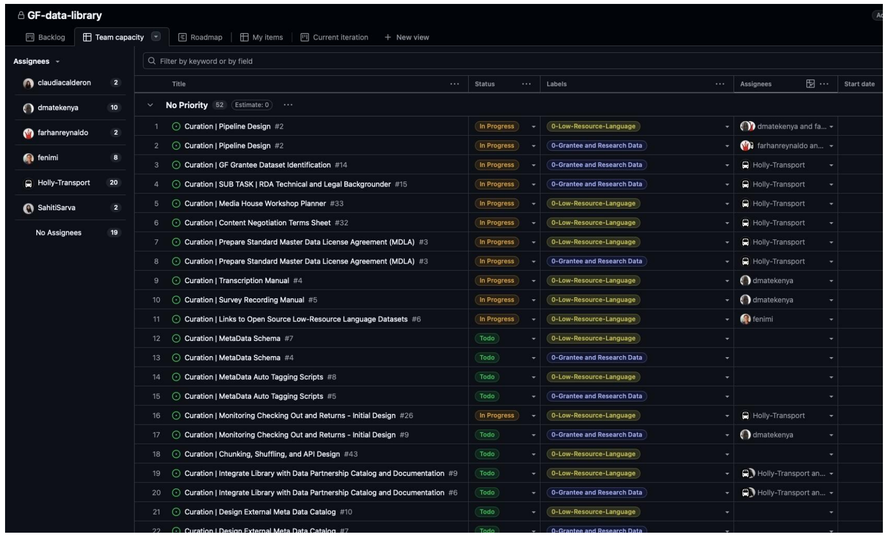

The Core Team

Holly Krambeck	Task Team Leader
Dunstan Matekenya	Data Scientist
Yoon Jee Kim	Legal Counsel
Keongmin Yoon	Legal Counsel
Aivin Solatorio	Data Scientist - Advisor
Sahiti Sarva	Data Scientist - DDP
Maria Sol Tadeo	Data Scientist - DDP
Claudia Calderon	Partnership Specialist - DDP
Evance Mathewe	Consultant – Malawi Pilot
Ifeoluwanimi Esther Ogbeba	Consultant – Data Science
Towera Moyo	Consultant – Data Science

Milestones

February 2025 – June 2026

	Q1-3 2025	Q4 2025	Q1 2026	Q2 2026	Q3 2026
Curation	• Data owner consultations • Legal terms sheet	• Legal agreements • Data collection • Metadata schema	• **Curation Documentation Completion**	• Continued data collection / curation	• Continued data collection / curation
Processing		• Processing documentation • Processing implementation	• Data privacy and warranty checks	• **Processing Documentation Completion**	• Continued data processing
Management	• LRL government endorsement	• Grantee dataset application procedure	• Grantee library launch • API Solution Developed	• LRL model 'leakage' verification tool • LRL LLM/ASR guidance • LRL Library launch	• **Library Management Documentation Completion**
Capacity Building & Scaling					• Training Course • Complete Toolkit

DEVELOPMENT DATA PARTNERSHIP

THANK YOU

Development Data Partnership